

Chapter 4. Using SQL with PostgreSQL

Table of Contents

[Introduction to psql](#)[Using Tables](#)[Adding Data with INSERT and COPY](#)[Retrieving Rows with SELECT](#)[Modifying Rows with UPDATE](#)[Removing Rows with DELETE](#)[Using Sub-Queries](#)[Using Views](#)[Further SQL Application](#)

In this chapter we continue to discuss SQL, this time with a practical focus. We'll address creating tables, populating tables with data, and managing that data via SQL statements.

Like most network-capable database systems, PostgreSQL fits into a client-server paradigm. The heart of PostgreSQL is the server backend, or the *postmaster* process. It is called a "backend" because it is not meant to directly interface with a user; rather, it can be connected to with a variety of clients.

When you start the PostgreSQL service, the *postmaster* process starts running in the background, listening to a specific TCP/IP port for connections from clients. Unless explicitly configured, *postmaster* will bind to, and listen on, port 5432.

There are several interfaces available through which clients may connect to the *postmaster* process. The examples in this book use *psql*, the most portable and readily accessible client distributed with PostgreSQL.

This chapter covers *psql* basics, how to create and use tables, and how to retrieve and manage data within those tables. It also addresses SQL sub-queries and views.

Introduction to psql

The *psql* client is a command-line client distributed with PostgreSQL. It is often called the *interactive monitor* or *interactive terminal*. With *psql*, you get a simple yet powerful tool with which you can directly interface with the PostgreSQL server, and thereby begin exploring SQL.

Starting psql

Before starting *psql*, be sure that you have either copied the *psql* binary into a path in your system `PATH` variable (e.g., `/usr/bin`), or that you have placed the PostgreSQL binary path (e.g., `/usr/local/pgsql/bin`) within your list of paths in your `PATH` environment variable (as shown in [Chapter 2](#)).

How you set the appropriate `PATH` variable will depend on your system shell. An example in either `bash` or `ksh` might read:

```
$ export PATH=$PATH:/usr/local/pgsql/bin
```

An example in either `cs`h or `tc`sh might read:

```
$ set path=($path /usr/local/pgsql/bin)
```

Example 4-1. Setting system path for `psql`

```
[user@host user]$ psql
bash: psql: command not found
[user@host user]$ echo $PATH
/bin:/usr/bin:/usr/local/bin:/usr/bin/X11:/usr/X11R6/bin
[user@host user]$ export PATH=$PATH:/usr/local/pgsql/bin
[user@host user]$ psql testdb
Welcome to psql, the PostgreSQL interactive terminal.
```

```
Type: \copyright for distribution terms
      \h for help with SQL commands
      \? for help on internal slash commands
      \g or terminate with semicolon to execute query
      \q to quit
```

```
testdb=#
```

Note that [Example 4-1](#) takes place within a bash shell.

Once you have appropriately set your `PATH` variable, you should be able to type *psql*, along with a database name, to start up the PostgreSQL interactive terminal.

Warning

Shell environment variables are erased after you have logged out. If you wish for your changes to the `PATH` variable to be retained upon logging in, you need to enter the appropriate `PATH` declaration into your shell-specific start-up scripts (e.g., `~/.bash_profile`).

Introduction to `psql` Syntax

Upon starting *psql*, you are greeted with a brief synopsis of four essential *psql slash commands*: `\h` for SQL help, `\?` for help on *psql*-specific commands, `\g` for executing queries and `\q` for actually exiting *psql* once you are done.

Every *psql*-specific command is prefixed by a backslash; hence the term "slash command" used earlier. For a complete list of slash commands and a brief description their functions, type `\?` into the *psql* command line, and press enter.

Example 4-2. Listing `psql` slash commands

```
booktown=# \?
\a                toggle between unaligned and aligned mode
\c[onnect] [dbname|- [user]]
                  connect to new database (currently 'booktown')
\C <title>        table title
\copy ...         perform SQL COPY with data stream to the client machine
\copyright        show PostgreSQL usage and distribution terms
\d <table>        describe table (or view, index, sequence)
\d{t|i|s|v}      list tables/indices/sequences/views
\d{p|S|l}        list permissions/system tables/objects
\da              list aggregates
\dd [object]     list comment for table, type, function, or operator
\df              list functions
\do              list operators
```

```

\dT          list data types
\e [file]    edit the current query buffer or [file] with external editor
\echo <text> write text to stdout
\encoding <encoding> set client encoding
\f <sep>     change field separator
\g [file]    send query to backend (and results in [file] or |pipe)
\h [cmd]     help on syntax of sql commands, * for all commands
\H          toggle HTML mode (currently off)
\i <file>    read and execute queries from <file>
\l          list all databases
\lo_export, \lo_import, \lo_list, \lo_unlink
            large object operations
\o [file]    send all query results to [file], or |pipe
\p          show the content of the current query buffer
\pset <opt>  set table output <opt> = {format|border|expanded|fieldsep|
            null|recordsep|tuples_only|title|tableattr|pager}
\q          quit psql
\qecho <text> write text to query output stream (see \o)
\r          reset (clear) the query buffer
\s [file]    print history or save it in [file]
\set <var> <value> set internal variable
\t          show only rows (currently off)
\T <tags>    HTML table tags
\unset <var> unset (delete) internal variable
\w <file>    write current query buffer to a <file>
\x          toggle expanded output (currently off)
\z          list table access permissions
\! [cmd]     shell escape or command

```

Executing Queries

Entering and executing queries within *psql* can be done two different ways. When using the client in interactive mode, the normal method is to directly enter queries into the prompt (i.e., standard input, or `stdin`). However, through the use of *psql*'s `\i` slash command, you can have *psql* read and interpret a file on your local filesystem as the query data.

Entering queries at the psql prompt

To enter queries directly into the prompt, open *psql* and make sure you are connected to the correct database (and logged in as the correct user). You will be presented with a prompt that, by default, is set to display the name of the database you are currently connected to. The prompt will look like this: *psql*:

```
testdb=#
```

To pass SQL statements to PostgreSQL, simply type them into the prompt. Anything you type (barring a slash command) will be queued until you terminate the query with a semicolon. This is the case even if you start a new line of type, thus allowing you to spread query statements across multiple lines. Examine [Example 4-3](#) to see how this is done.

Example 4-3. Entering statements into psql

```
testdb=# SELECT * FROM employees
testdb=#         WHERE firstname = 'Michael';
```

The query entered in [Example 4-3](#) will return a table that consists of all employees whose first name is Michael. The query could be broken up over multiple lines to improve readability, and *psql* would not send it to the backend until the terminating semicolon was sent. The prompt will show the end-character of a previous line if the character requires a closing character, such as a parenthesis or a quote (this is not

shown in the example). If you were to issue a `CREATE TABLE` command to start a statement, and then hit enter to begin a new line for readability purposes, you would see a prompt similar to the one displayed in [Example 4-4](#).

Example 4-4. Leaving end-characters open

```
testdb=# CREATE TABLE employees (  
testdb(
```

At this point you could continue the statement. The *psql* prompt is informing you of the open parenthesis by inserting an open parenthesis symbol into the prompt.

Editing the query buffer

Use the `\e` command to edit the current query buffer with the editor that your `EDITOR` environment variable is set to. Doing so can be very useful when entering queries and statements in *psql*, as you can easily view and modify all lines of your query or statement before it is committed. [Example 4-5](#) shows how to set the `EDITOR` variable. The *vi* editor will be used if `EDITOR` is not set.

Example 4-5. Setting the EDITOR variable

```
$ set EDITOR='joe'  
$ export EDITOR
```

You can also use this command to save your current buffer as a file. Issue the `\e` command to enter editing mode. This will open your editor and load the buffer as if it were a file. Complete whatever work you wish to do with the buffer, then use your editor's save function to save the buffer and return to *psql*. To save the query as a normal file, use your editor's save-as function and save it as a file other than the *.tmp* created by `\e`.

[Prev](#)[Tables in PostgreSQL](#)[Home](#)[Up](#)[Next](#)[Using Tables](#)